


Motivation using RNN




Motivation using RNN

Analyzing temporal dependencies



Improved decisions




Decision on
sequence of
observations



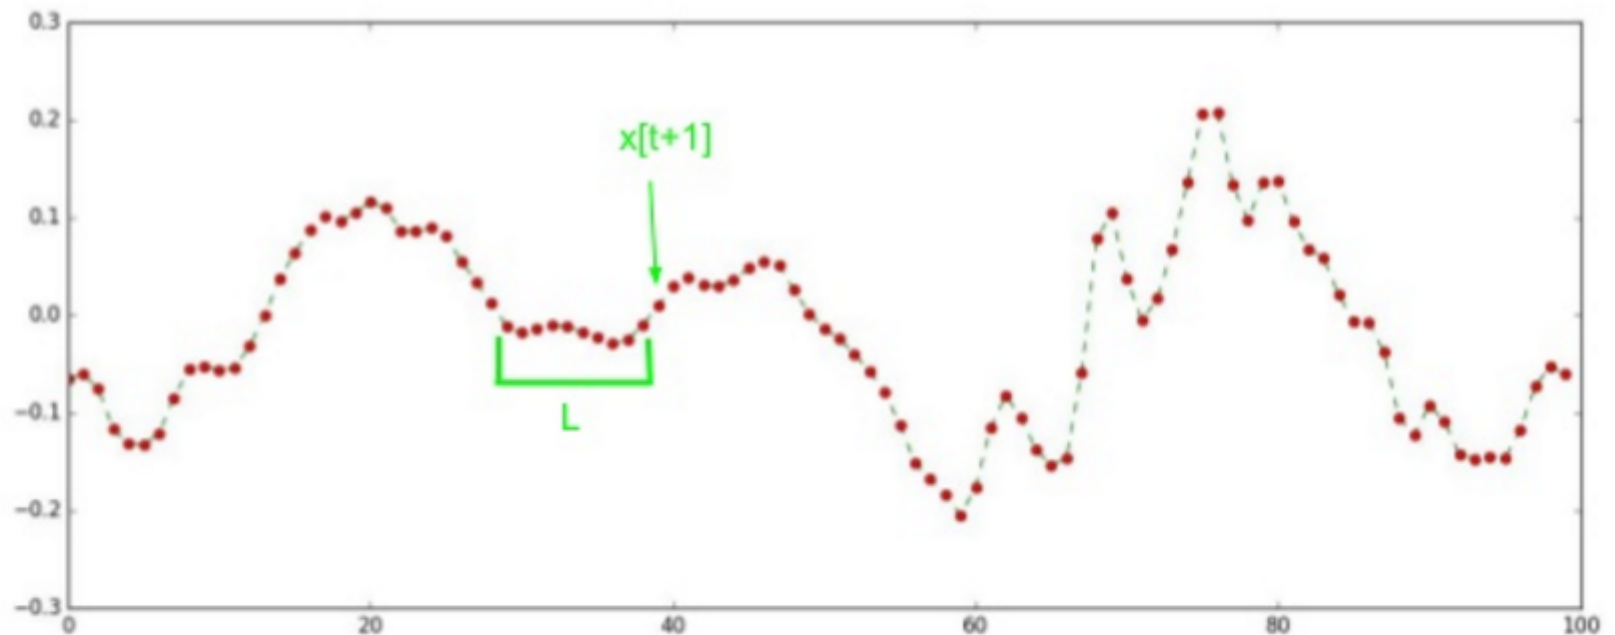
From FF to RNN

If we have a sequence of samples...




predict sample $x[t+1]$ knowing previous values $\{x[t], x[t-1], x[t-2], \dots, x[t-\tau]\}$

From FF to RNN

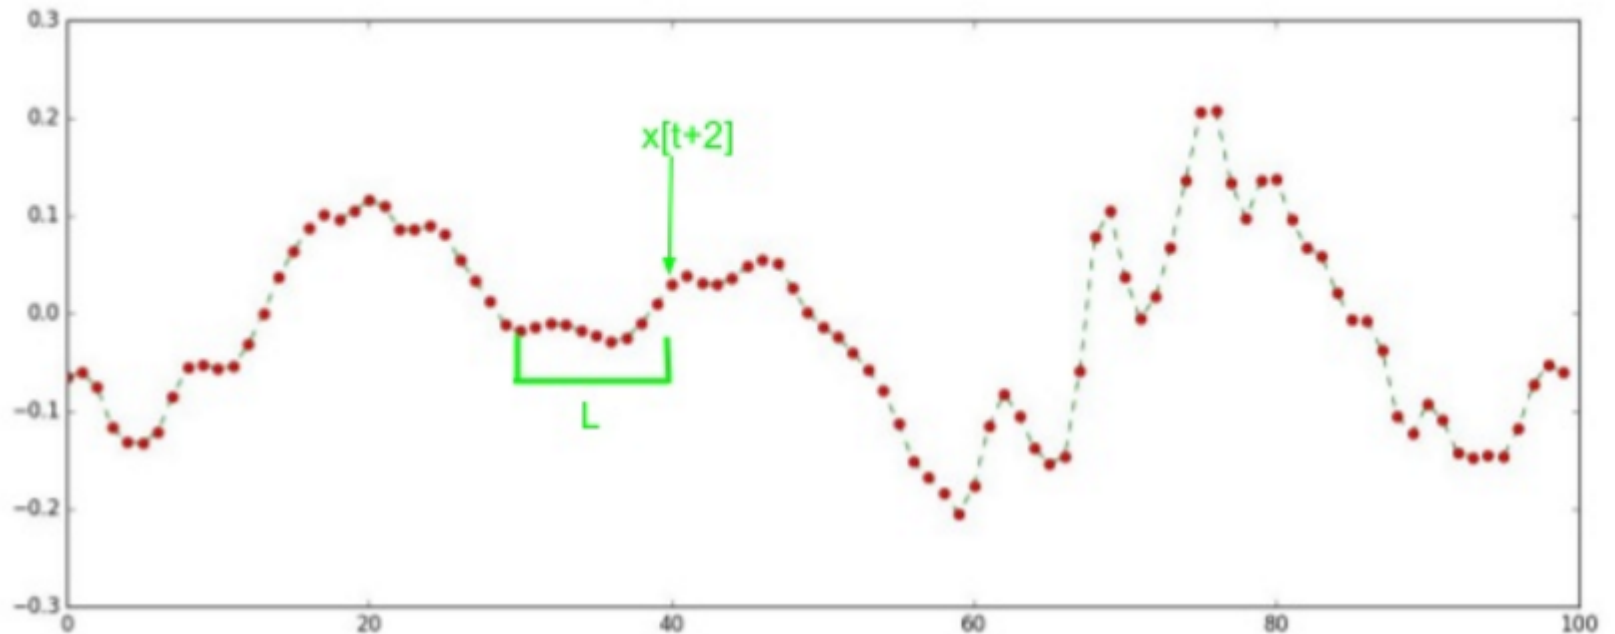


Feed Forward approach:

- static window of size L
- slide the window time-step wise




From FF to RNN




Feed Forward approach:

- static window of size L
- slide the window time-step wise




From FF to RNN



Feed Forward approach:


- static window of size L
- slide the window time-step wise



From FF to RNN

Problems for the feed forward + static window approach:

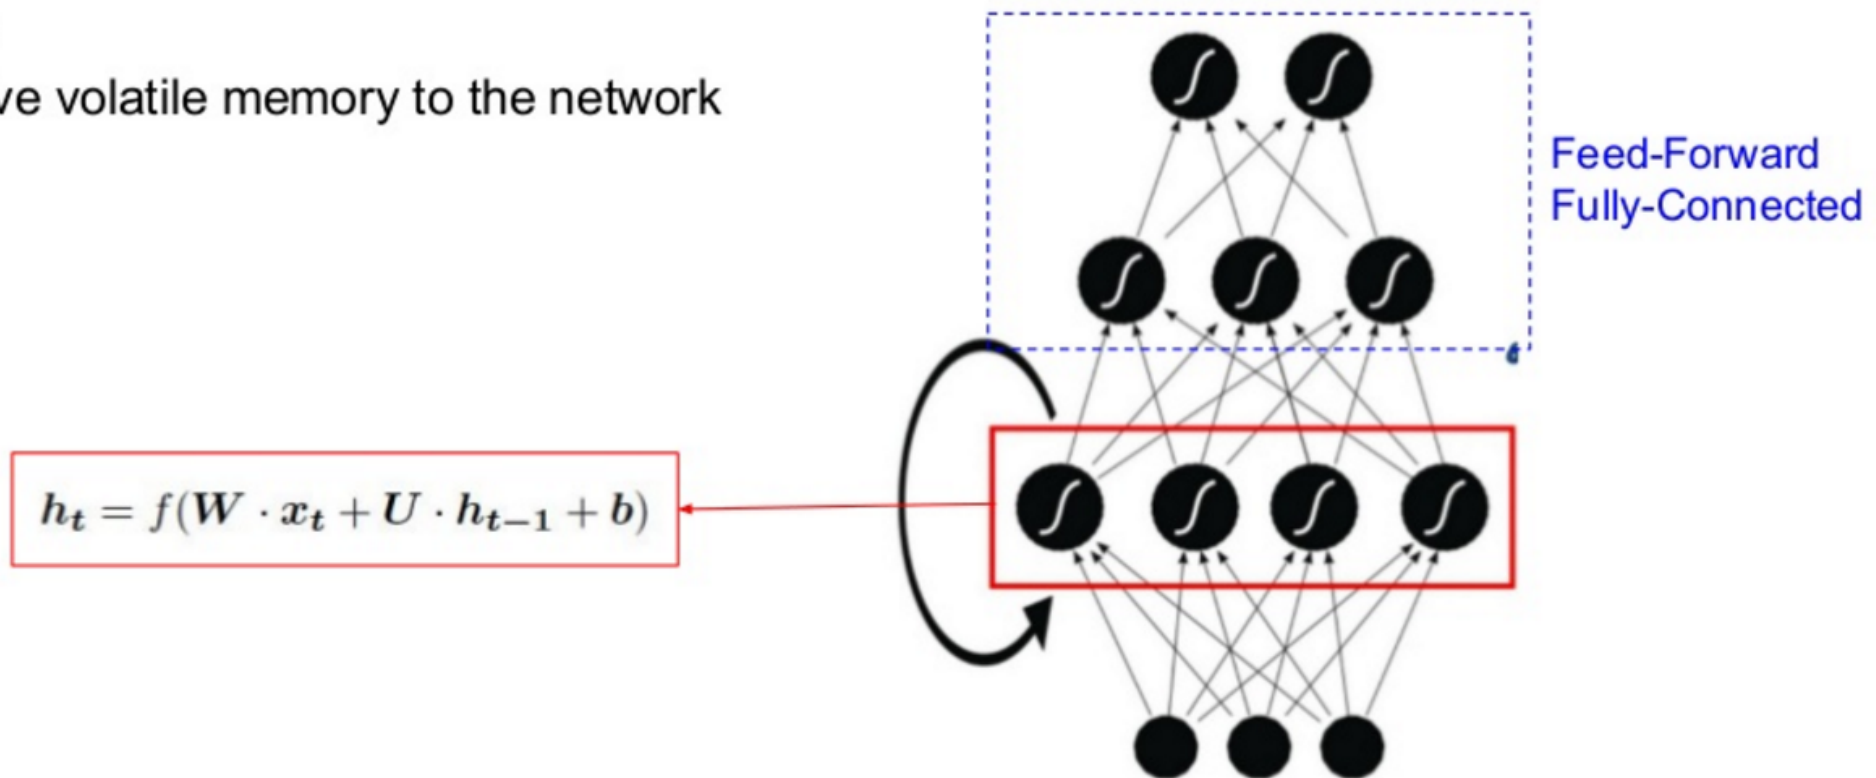
- What's the matter increasing L ? → Fast growth of num of parameters!
- Decisions are independent between time-steps!
 - The network doesn't care about what happened at previous time-step, only present window matters → doesn't look good
- Cumbersome padding when there are not enough samples to fill L size
 - Can't work with variable sequence lengths



From FF to RNN


Solution: Build specific connections capturing the temporal evolution → **Shared weights in time**

- Give volatile memory to the network



RNN

- An MLP can only map from input to output vectors, whereas an RNN can, in principle, map from the entire history of previous inputs to each output.




Multi-layer
Perceptron

Output Layer

Hidden Layers


Input Layer




Recurrent Network

RNN offer a lot of flexibility


one to one




one to many




many to one



many to many




many to many




Vanilla Neural Networks

RNN offer a lot of flexibility


one to one




one to many




many to one



many to many




many to many




↖ e.g. **Image Captioning**
image -> sequence of words

RNN offer a lot of flexibility


one to one




one to many




many to one



many to many




many to many




e.g. **Sentiment Classification**
sequence of words -> sentiment

RNN offer a lot of flexibility


one to one




one to many




many to one



many to many




many to many




↖ e.g. **Machine Translation**
seq of words -> seq of words

RNN offer a lot of flexibility


one to one




one to many




many to one



many to many



many to many




e.g. Video classification on frame level




The problem of long-term dependencies

- (Vanilla) RNNs connect previous information to present task:
- - enough for predicting the next word for “the clouds are in the *sky*”




- - may not be enough when more context is needed
- “I grew up in France... I speak fluent *French*.”



The problem of vanishing gradients


- In a traditional recurrent neural network, during the gradient backpropagation phase, the gradient signal can end up being multiplied a large number of times
- If the gradients are large
 - Exploding gradients, learning diverges
 - **Solution: Clip the gradients to a certain max value.**
- If the gradients are small
 - Vanishing gradients, learning very slow or stops
 - **Solution: introducing memory via LSTM, GRU, etc.**



Long Short Term Memory (LSTM) [Hochreiter & Schmidhuber (1997)]

- Learning long-term dependencies is difficult with simple RNNs, unstable training due to vanishing gradients problem [Hochreiter, 1991]
- Limited capability (5-10 time steps) to model long-term dependencies
- LSTM RNN architecture designed to address these problems [Hochreiter and Schmidhuber, 1997]
- LSTM memory block: memory cell storing temporal state of network and 3 multiplicative units (gates) controlling the flow of information

Long Short Term Memory (LSTM) [Hochreiter & Schmidhuber (1997)]



Long Short Term Memory (LSTM) [Hochreiter & Schmidhuber (1997)]

Key difference to RNN: a memory cell c which is controlled by three gates:

- ▶ input gate i :

$$i_t = \sigma(W_i x_t + U_i h_{t-1} + b_i) \quad (7)$$

- ▶ output gate o :

$$o_t = \sigma(W_o x_t + U_o h_{t-1} + b_o) \quad (8)$$

- ▶ forget gate f :

$$f_t = \sigma(W_f x_t + U_f h_{t-1} + b_f) \quad (9)$$

- ▶ update of memory cell and hidden state:

$$\widetilde{C}_t = \tanh(W_c x_t + U_c h_{t-1} + b_c) \quad (10)$$


$$C_t = f_t \otimes C_{t-1} + i_t \otimes \widetilde{C}_t \quad (11)$$

$$h_t = o_t \otimes \tanh(C_t) \quad (12)$$

Long Short Term Memory (LSTM) [Hochreiter & Schmidhuber (1997)]

- LSTM performs better than RNN for learning context-free and context-sensitive languages [Gers and Schmidhuber, 2001]
- Bidirectional LSTM for phonetic labeling of acoustic frames on the TIMIT [Graves and Schmidhuber, 2005]
- Online and offline handwriting recognition with bidirectional LSTM better than HMM-based system [Graves et al., 2009]
- Deep LSTM - stack of multiple LSTM layers - combined with CTC and RNN transducer predicting phone sequences gets state of the art results on TIMIT [Graves et al., 2013]

Gated Recurrent Unit



Gated Recurrent Unit

Instead of using three gates, only two gates is employed:

- update gate z :

$$z_t = \sigma(W_z x_t + U_z h_{t-1}) \quad (13)$$

- reset gate r :

$$r_t = \sigma(W_r x_t + U_r h_{t-1}) \quad (14)$$

- state update:

$$\tilde{h}_t = \tanh(W_c x_t + U(r_t \otimes h_{t-1})) \quad (15)$$

$$h_t = (1 - z_t) \otimes h_{t-1} + z_t \otimes \tilde{h}_t \quad (16)$$

Multivariate Forecasting

Beijing Air Quality Dataset

- No: row number
- year: year of data in this row
- month: month of data in this row
- day: day of data in this row
- hour: hour of data in this row
- pm2.5: PM2.5 concentration
- DEWP: Dew Point
- TEMP: Temperature
- PRES: Pressure
- cbwd: Combined wind direction
- lws: Cumulated wind speed
- ls: Cumulated hours of snow
- lr: Cumulated hours of rain

Multivariate Forecasting

